

Data anomaly: inconsistency in DB as result of an update, insertion, and/or deletion

Logical independence: data is independent of the table structure, i.e. unaffected when a col is added or a col is moved

Physical independence: the data is independent from its physical container and unaffected by changes

Functional dependencies: partial and transitive dependencies

Locking: controls access that users have to the DB

Exclusivity: maintain data consistency when same data updated simultaneously by multiple transactions

Granularity: defines how fine the control of the lock is

Lost updates: 2 diff transactions try to update same col on same row within DB at same time, result of one is lost

Uncommitted data: 1 transaction updates data, but hasn't committed it permanently to the DB, b/c failure, transaction is rolled back and data item is returned to previous value

Inconsistent retrievals: transaction calculates summary function over set of data while other transactions update data

Concurrency: ability to allow multiple users to affect multiple transactions within a DB (multiple users use DB at once)

Concurrency control: manages simultaneous access to DB, prevents 2 users from editing same record at same time

Transaction: logical unit independently executed for data retrieval or updates, group of tasks (like a query for ex)

Transaction logfil: records all transactions and DB modifications made by each transaction (good for troubleshooting)

Snapshots: backup of DB

ACID: Atomicity, Consistency, Isolation, Durability

Atomicity: transaction treated as atomic unit, either all operations executed or none, no partial transactions

DB Consistency: DB remain in consistent state after any transaction, transaction should not have effect on data

Isolation: all transactions being executed simultaneously are treated independently

Durability: durable enough to hold latest updates even if system fails or restarts, committed transactions not lost

COMMIT: permanently save changes done in transaction in tables/DBs. DB cannot regain previous state after this

ROLLBACK: undo transactions that have not been saved in DB, undoes all changes since last COMMIT

Performance considerations for locking, exclusivity, granularity:

More difficult to access all of the data but safer, lower performance the more restrictive it is

Lock Granularities for accessibility and performance cost

DB: 1 person controls all data in DB

Table: select tables can be locked/accessed

Row: can only access certain rows

Column: can only access certain cells of all rows

As you go down, requires more resources and is harder to implement

SQL

Query Syntax: *select vals from table1 t1 join table2 t2 on t1.id = t2.id where conditions order by val asc/desc;*

Linking Table: *insert into linkingTable (varNames) select id1, id2 from table1 join table2 on ... where ...;* (make combos of keys)

Create: *create table tableName (varName varType primary key auto_increment, varName2 varType2 not null, ...);*

Create Combo Key: create table w/o PK, then add: *primary key(pkVal1, pkVal2)*

Foreign Key: add val as normal val, then add *foreign key (valName) references referenceTable(referenceKey)*

Insert: *insert into tableName (fieldsInserted) values (valsInOrder);*

Or: *insert into tableName (fieldsInserted) select val1, val2, ... from tables;* (instead of values... , do a query)

Codd's Rules: **1)** all info in RDB is logically represented as col vals in rows in tables. **2)** all vals in table must be accessible thru combo of table name, PK val, col name. **3)** nulls represented and treated in systematic way, indep of data type. **4)** metadata stored and managed as ordinary data, in tables in DB, data is available to allowed users using standard DB lang. **5)** RDB can have many langs, but must support 1 well defined declarative lang w/ data defn, view defn, data manip, integrity constraints, authorization, transaction management. **6)** any view that is theoretically updatable must be updatable thru the system. **7)** DB supports set level insert, update, delete. **8)** app programs/ad hoc facilities logically unaffected when physical access methods or storage structures are changed. **9)** app progs and ad hoc facilities logically unaffected when change made to table structures that keep original table vals (change order of cols or insert col). **10)** relational integrity constraints must be definable in relational lang & stored in system catalog, not at app level. **11)** end user and app progs unaware and unaffected by data location. **12)** if system supports low level access to data, must not be way to bypass integrity rules of DB.